

AutoRec: Autoencoders Meet Collaborative Filtering

Introduction

Problem: Given U users and their preferences (star ratings) about some of the given I items. The task is to find preferences about the remaining items.

This can also be viewed as a problem of recommendation where we want to recommend items to users depending on the preferences which they have given.

Collaborative Filtering

Collaborative Filtering (Main Idea) - If a person A has the same opinion as a person B on an issue, A is more likely to have B 's opinion on a different issue than that of a randomly chosen person.

Mathematical Representation -
$$r_{u,i} = \bar{r}_u + k \sum_{u' \in U} \text{simil}(u, u') (r_{u',i} - \bar{r}_{u'})$$

where \bar{r}_u is average rating given by user u and k is defined as $k = 1 / \sum_{u' \in U} |\text{simil}(u, u')|$

where $\text{simil}(u, v)$ is similarity measure which can be cosine similarity or any other similarity measure. This is basic collaborative filtering approach.

Problem setting

- Rating-based collaborative filtering:
 - m users
 - n items
 - Partially observed rating matrix R
 - Each item - partially observed vector $\mathbf{r}^{(i)} = (R_{1i}, \dots, R_{mi})$
- Task: Design an item-based autoencoder which can:
 - take as input each partially observed $\mathbf{r}^{(i)}$
 - project it into a low-dimensional latent space
 - and then reconstruct $\mathbf{r}^{(i)}$ in the output space


Formulation

$$\min_{\theta} \sum_{\mathbf{r} \in \mathbf{S}} \|\mathbf{r} - h(\mathbf{r}; \theta)\|_2^2,$$

where $h(\mathbf{r}; \theta)$ is the *reconstruction* of input $\mathbf{r} \in \mathbb{R}^d$,

$$h(\mathbf{r}; \theta) = f(\mathbf{W} \cdot g(\mathbf{V}\mathbf{r} + \boldsymbol{\mu}) + \mathbf{b})$$

Model



of parameters: $2mk + m + k$

Objection function with regularization:

$$\min_{\theta} \sum_{i=1}^n \|\mathbf{r}^{(i)} - h(\mathbf{r}^{(i)}; \theta)\|_{\mathcal{O}}^2 + \frac{\lambda}{2} \cdot (\|\mathbf{W}\|_F^2 + \|\mathbf{V}\|_F^2)$$

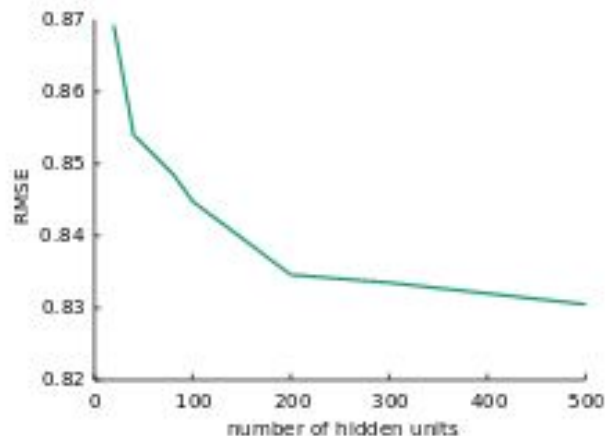
Using the model:

$$\hat{R}_{ui} = (h(\mathbf{r}^{(i)}; \hat{\theta}))_u$$

Experimental Results

1. The dataset used for experimentation was movielens dataset which consisted of movie ratings of 1 million users.
2. The RMSE values obtained for different methods are given below.
3. Performance steadily increases with the number of hidden units, but with diminishing returns.

	ML-1M	ML-10M
U-RBM	0.881	0.823
I-RBM	0.854	0.825
U-AutoRec	0.874	0.867
I-AutoRec	0.831	0.782



Second order optimization methods:

A very brief overview

Newton's method

$$\mathbf{x}_{n+1} = \mathbf{x}_n - \gamma[\mathbf{H}f(\mathbf{x}_n)]^{-1} \nabla f(\mathbf{x}_n)$$

- Computing the inverse when the data is high dimensional could be very expensive
 - Alternative 1: Pose the problem of finding updates as a system of linear equations
 - Alternative 2: Compute the Hessian (or its inverse directly) from changes in the gradient (Quasi-Newton methods).

BFGS

- Broyden–Fletcher–Goldfarb–Shanno (BFGS) algorithm
- Among the most popular Quasi Newton methods
- It stores the approximation to the Hessian matrix.
- If the dimension of the x vector is n , that takes $O(n^2)$ memory.

L-BFGS

- Limited-memory (L) Broyden–Fletcher–Goldfarb–Shanno (BFGS) algorithm
- L-BFGS is a BFGS version that does not require to explicitly store the approximation to the Hessian
 - So, it needs less memory